

Name: DHARMANA GNANA SAI

Reg No: 21BCE7400

Assignment-4

• Data Preprocessing.

- o Import the Libraries.
- o Importing the dataset.
- o Checking for Null Values.
- o Data Visualization.
- o Outlier Detection
- o Splitting Dependent and Independent variables
- o- Encoding
- o Feature Scaling.
- o Splitting Data into Train and Test.

Data Collection. o Collect the dataset or Create the dataset • Data Preprocessing. o Import the Libraries. o Importing the dataset. o Checking for Null Values. o Data Visualization. o Outlier Detection o Splitting Dependent and Independent variables o- Encoding o Feature Scaling. o Splitting Data into Train and Test. • Model Building o Import the model building Libraries o Initializing the model o Training and testing the model o Evaluation of Model o Save the Model


```
In [1]: #Import the Libraries.  
import numpy as np  
import pandas as pd  
import seaborn as sns  
import matplotlib.pyplot as plt
```

```
In [2]: #Importing the dataset.  
df = pd.read_csv('Employee-Attrition.csv')  
df
```

Out[2]:

	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Ed
0	41	Yes	Travel_Rarely	1102	Sales		1
1	49	No	Travel_Frequently	279	Research & Development		8
2	37	Yes	Travel_Rarely	1373	Research & Development		2
3	33	No	Travel_Frequently	1392	Research & Development		3
4	27	No	Travel_Rarely	591	Research & Development		2
...
1465	36	No	Travel_Frequently	884	Research & Development		23
1466	39	No	Travel_Rarely	613	Research & Development		6
1467	27	No	Travel_Rarely	155	Research & Development		4
1468	49	No	Travel_Frequently	1023	Sales		2
1469	34	No	Travel_Rarely	628	Research & Development		8

1470 rows × 35 columns




In [3]: df.head()

Out[3]:

	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Educational
0	41	Yes	Travel_Rarely	1102	Sales		1
1	49	No	Travel_Frequently	279	Research & Development		8
2	37	Yes	Travel_Rarely	1373	Research & Development		2
3	33	No	Travel_Frequently	1392	Research & Development		3
4	27	No	Travel_Rarely	591	Research & Development		2

5 rows × 35 columns



In [4]: df.shape

Out[4]: (1470, 35)

```
In [5]: df.DailyRate.value_counts()
```

```
Out[5]: 691      6
        408      5
        530      5
        1329     5
        1082     5
        ..
        650      1
        279      1
        316      1
        314      1
        628      1
        Name: DailyRate, Length: 886, dtype: int64
```

```
In [6]: df.info()
```


```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1470 entries, 0 to 1469
Data columns (total 35 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Age                                   1470 non-null   int64
1   Attrition                           1470 non-null   object
2   BusinessTravel                       1470 non-null   object
3   DailyRate                           1470 non-null   int64
4   Department                           1470 non-null   object
5   DistanceFromHome                    1470 non-null   int64
6   Education                           1470 non-null   int64
7   EducationField                       1470 non-null   object
8   EmployeeCount                       1470 non-null   int64
9   EmployeeNumber                      1470 non-null   int64
10  EnvironmentSatisfaction              1470 non-null   int64
11  Gender                               1470 non-null   object
12  HourlyRate                           1470 non-null   int64
13  JobInvolvement                       1470 non-null   int64
14  JobLevel                             1470 non-null   int64
15  JobRole                              1470 non-null   object
16  JobSatisfaction                      1470 non-null   int64
17  MaritalStatus                       1470 non-null   object
18  MonthlyIncome                       1470 non-null   int64
19  MonthlyRate                         1470 non-null   int64
20  NumCompaniesWorked                  1470 non-null   int64
21  Over18                              1470 non-null   object
22  OverTime                             1470 non-null   object
23  PercentSalaryHike                   1470 non-null   int64
24  PerformanceRating                   1470 non-null   int64
25  RelationshipSatisfaction             1470 non-null   int64
26  StandardHours                       1470 non-null   int64
27  StockOptionLevel                    1470 non-null   int64
28  TotalWorkingYears                   1470 non-null   int64
29  TrainingTimesLastYear               1470 non-null   int64
30  WorkLifeBalance                     1470 non-null   int64
31  YearsAtCompany                      1470 non-null   int64
32  YearsInCurrentRole                  1470 non-null   int64
33  YearsSinceLastPromotion              1470 non-null   int64
34  YearsWithCurrManager                 1470 non-null   int64
dtypes: int64(26), object(9)
memory usage: 402.1+ KB
```

```
In [7]: df.describe()
```

Out[7]:

	Age	DailyRate	DistanceFromHome	Education	EmployeeCount	Em
count	1470.000000	1470.000000	1470.000000	1470.000000	1470.0	
mean	36.923810	802.485714	9.192517	2.912925	1.0	
std	9.135373	403.509100	8.106864	1.024165	0.0	
min	18.000000	102.000000	1.000000	1.000000	1.0	
25%	30.000000	465.000000	2.000000	2.000000	1.0	
50%	36.000000	802.000000	7.000000	3.000000	1.0	
75%	43.000000	1157.000000	14.000000	4.000000	1.0	
max	60.000000	1499.000000	29.000000	5.000000	1.0	

8 rows × 26 columns



```
In [8]: #Checking for Null Values.  
df.isnull().any()
```

```
Out[8]: Age False
Attrition False
BusinessTravel False
DailyRate False
Department False
DistanceFromHome False
Education False
EducationField False
EmployeeCount False
EmployeeNumber False
EnvironmentSatisfaction False
Gender False
HourlyRate False
JobInvolvement False
JobLevel False
JobRole False
JobSatisfaction False
MaritalStatus False
MonthlyIncome False
MonthlyRate False
NumCompaniesWorked False
Over18 False
OverTime False
PercentSalaryHike False
PerformanceRating False
RelationshipSatisfaction False
StandardHours False
StockOptionLevel False
TotalWorkingYears False
TrainingTimesLastYear False
WorkLifeBalance False
YearsAtCompany False
YearsInCurrentRole False
YearsSinceLastPromotion False
YearsWithCurrManager False
dtype: bool
```

```
In [9]: df.isnull().sum()
```

```
Out[9]: Age 0
Attrition 0
BusinessTravel 0
DailyRate 0
Department 0
DistanceFromHome 0
Education 0
EducationField 0
EmployeeCount 0
EmployeeNumber 0
EnvironmentSatisfaction 0
Gender 0
HourlyRate 0
JobInvolvement 0
JobLevel 0
JobRole 0
JobSatisfaction 0
MaritalStatus 0
MonthlyIncome 0
MonthlyRate 0
NumCompaniesWorked 0
Over18 0
OverTime 0
PercentSalaryHike 0
PerformanceRating 0
RelationshipSatisfaction 0
StandardHours 0
StockOptionLevel 0
TotalWorkingYears 0
TrainingTimesLastYear 0
WorkLifeBalance 0
YearsAtCompany 0
YearsInCurrentRole 0
YearsSinceLastPromotion 0
YearsWithCurrManager 0
dtype: int64
```

```
In [3]: #Data Visualization.
sns.distplot(df["Age"])
```

C:\Users\dharm\AppData\Local\Temp\ipykernel_14632\2400079689.py:2: UserWarning:


`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(df["Age"])
```

```
Out[3]: <Axes: xlabel='Age', ylabel='Density'>
```




```
In [11]: df.corr()
```

C:\Users\chatu\AppData\Local\Temp\ipykernel_8416\1134722465.py:1: FutureWarning:
The default value of numeric_only in DataFrame.corr is deprecated. In a future ve
rsion, it will default to False. Select only valid columns or specify the value o
f numeric_only to silence this warning.
df.corr()

Out[11]:

	Age	DailyRate	DistanceFromHome	Education	Employee
Age	1.000000	0.010661	-0.001686	0.208034	
DailyRate	0.010661	1.000000	-0.004985	-0.016806	
DistanceFromHome	-0.001686	-0.004985	1.000000	0.021042	
Education	0.208034	-0.016806	0.021042	1.000000	
EmployeeCount	NaN	NaN	NaN	NaN	
EmployeeNumber	-0.010145	-0.050990	0.032916	0.042070	
EnvironmentSatisfaction	0.010146	0.018355	-0.016075	-0.027128	
HourlyRate	0.024287	0.023381	0.031131	0.016775	
JobInvolvement	0.029820	0.046135	0.008783	0.042438	
JobLevel	0.509604	0.002966	0.005303	0.101589	
JobSatisfaction	-0.004892	0.030571	-0.003669	-0.011296	
MonthlyIncome	0.497855	0.007707	-0.017014	0.094961	
MonthlyRate	0.028051	-0.032182	0.027473	-0.026084	
NumCompaniesWorked	0.299635	0.038153	-0.029251	0.126317	
PercentSalaryHike	0.003634	0.022704	0.040235	-0.011111	
PerformanceRating	0.001904	0.000473	0.027110	-0.024539	
RelationshipSatisfaction	0.053535	0.007846	0.006557	-0.009118	
StandardHours	NaN	NaN	NaN	NaN	
StockOptionLevel	0.037510	0.042143	0.044872	0.018422	
TotalWorkingYears	0.680381	0.014515	0.004628	0.148280	
TrainingTimesLastYear	-0.019621	0.002453	-0.036942	-0.025100	
WorkLifeBalance	-0.021490	-0.037848	-0.026556	0.009819	
YearsAtCompany	0.311309	-0.034055	0.009508	0.069114	
YearsInCurrentRole	0.212901	0.009932	0.018845	0.060236	
YearsSinceLastPromotion	0.216513	-0.033229	0.010029	0.054254	
YearsWithCurrManager	0.202089	-0.026363	0.014406	0.069065	

26 rows × 26 columns




In [12]: df.head()

Out[12]:

	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Educational
0	41	Yes	Travel_Rarely	1102	Sales		1
1	49	No	Travel_Frequently	279	Research & Development		8
2	37	Yes	Travel_Rarely	1373	Research & Development		2
3	33	No	Travel_Frequently	1392	Research & Development		3
4	27	No	Travel_Rarely	591	Research & Development		2

5 rows × 35 columns

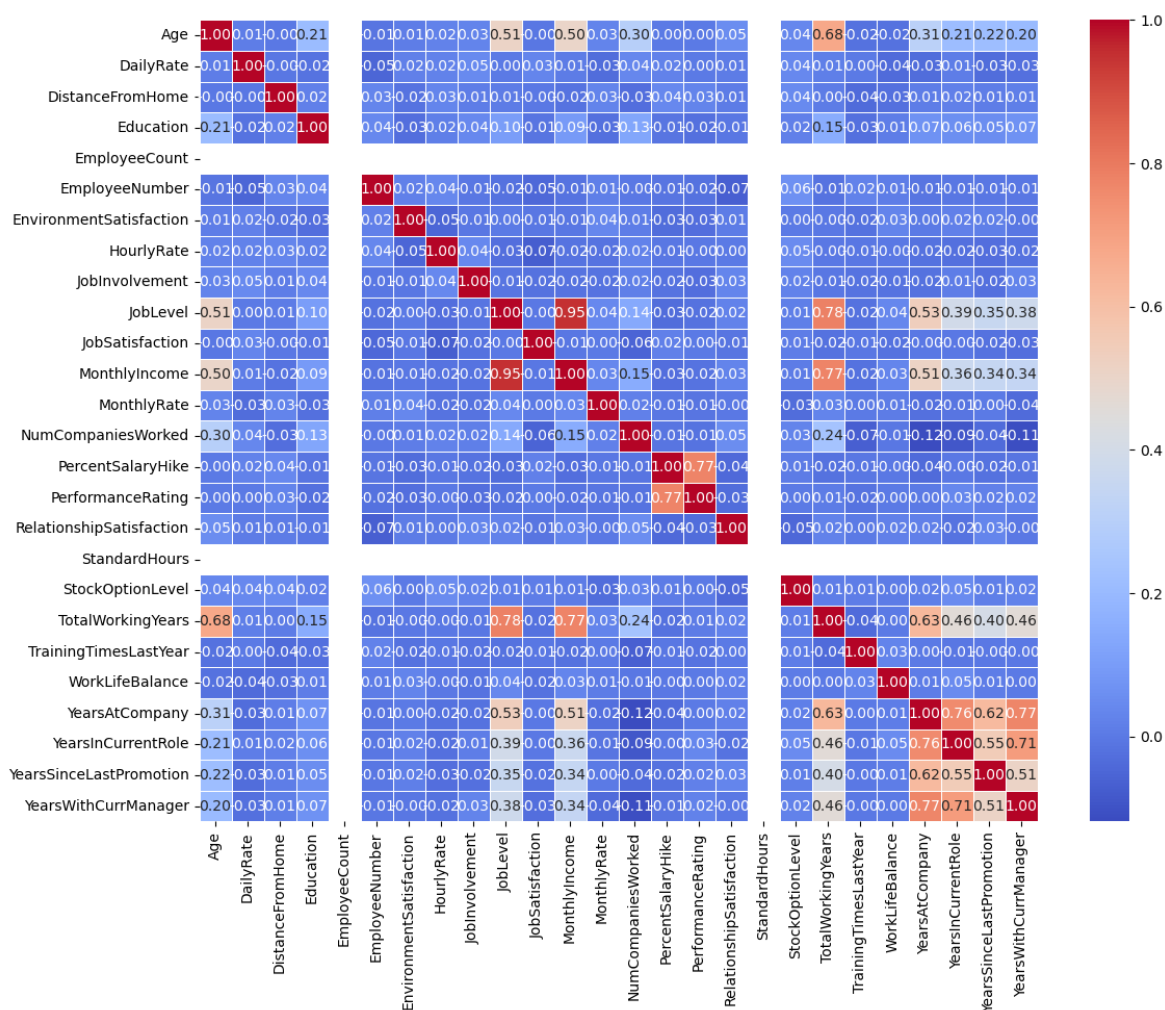


In [13]: `plt.figure(figsize=(13, 10))`
`sns.heatmap(df.corr(),annot=True,cmap='coolwarm', fmt='.2f', linewidths=0.5)`

C:\Users\chatu\AppData\Local\Temp\ipykernel_8416\195898717.py:2: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric_only to silence this warning.

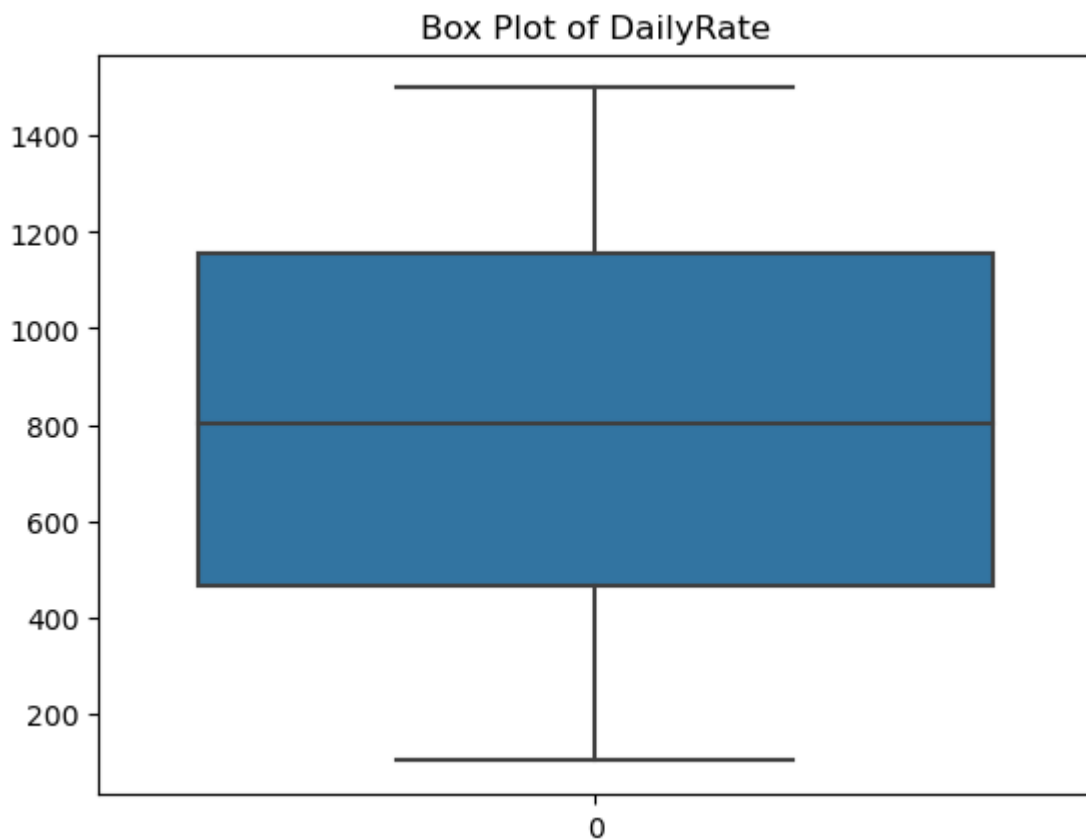
`sns.heatmap(df.corr(),annot=True,cmap='coolwarm', fmt='.2f', linewidths=0.5)`

Out[13]: <Axes: >




```
In [14]: plt.title("Box Plot of DailyRate")  
sns.boxplot(df.DailyRate)
```

```
Out[14]: <Axes: title={'center': 'Box Plot of DailyRate'}>
```



```
In [17]: plt.title("Box Plot of Age")  
sns.boxplot(df.Age)
```

```
Out[17]: <Axes: title={'center': 'Box Plot of Age'}>
```




In [18]: `df.head()`

Out[18]:

	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Educational
0	41	Yes	Travel_Rarely	1102	Sales	1	
1	49	No	Travel_Frequently	279	Research & Development	8	
2	37	Yes	Travel_Rarely	1373	Research & Development	2	
3	33	No	Travel_Frequently	1392	Research & Development	3	
4	27	No	Travel_Rarely	591	Research & Development	2	

5 rows × 35 columns



In [21]: `#Splitting Dependent and Independent variables`
`x=df.iloc[:,1:7]`
`x.head()`

Out[21]:

	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Education
0	Yes	Travel_Rarely	1102	Sales	1	2
1	No	Travel_Frequently	279	Research & Development	8	1
2	Yes	Travel_Rarely	1373	Research & Development	2	2
3	No	Travel_Frequently	1392	Research & Development	3	4
4	No	Travel_Rarely	591	Research & Development	2	1

```
In [22]: y=df.DistanceFromHome
y.head()
```

```
Out[22]: 0    1
1    8
2    2
3    3
4    2
Name: DistanceFromHome, dtype: int64
```

```
In [23]: #Label encoding
from sklearn.preprocessing import LabelEncoder
le=LabelEncoder()
x['DailyRate']=le.fit_transform(x['DailyRate'])
x.head()
```

Out[23]:

	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Education
0	Yes	Travel_Rarely	624	Sales	1	2
1	No	Travel_Frequently	113	Research & Development	8	1
2	Yes	Travel_Rarely	805	Research & Development	2	2
3	No	Travel_Frequently	820	Research & Development	3	4
4	No	Travel_Rarely	312	Research & Development	2	1

```
In [24]: #feature scaling
non_numeric_columns = x.select_dtypes(exclude=['number']).columns
# Drop non-numeric columns
x_numeric = x.drop(columns=non_numeric_columns)
x_encoded = pd.get_dummies(x, columns=non_numeric_columns)

from sklearn.preprocessing import MinMaxScaler
ms=MinMaxScaler()
x_scaled = pd.DataFrame(ms.fit_transform(x_numeric), columns=x_numeric.columns)
```

```
In [25]: x_scaled
```

Out[25]:

	DailyRate	DistanceFromHome	Education
0	0.705085	0.000000	0.25
1	0.127684	0.250000	0.00
2	0.909605	0.035714	0.25
3	0.926554	0.071429	0.75
4	0.352542	0.035714	0.00
...
1465	0.558192	0.785714	0.25
1466	0.369492	0.178571	0.00
1467	0.044068	0.107143	0.50
1468	0.654237	0.035714	0.50
1469	0.379661	0.250000	0.50

1470 rows × 3 columns

```
In [26]: #Splitting Data into Train and Test.
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x_scaled,y,test_size=0.2,random_s
```

```
In [27]: x_train.shape,x_test.shape,y_train.shape,y_test.shape
```

```
Out[27]: ((1176, 3), (294, 3), (1176,), (294,))
```

```
In [28]: x_train.head()
```

Out[28]:

	DailyRate	DistanceFromHome	Education
1374	0.364972	0.714286	0.50
1092	0.605650	0.964286	0.50
768	0.141243	0.892857	0.50
569	0.954802	0.250000	0.75
911	0.358192	0.821429	0.00

• Model Building

- o Import the model building Libraries
- o Initializing the model
- o Training and testing the model
- o Evaluation of Model
- o Save the Model

```
In [29]: from sklearn.tree import DecisionTreeClassifier
```

```
dtc=DecisionTreeClassifier()
```

```
In [30]: dtc.fit(x_train,y_train)
```

```
Out[30]: ▼ DecisionTreeClassifier
DecisionTreeClassifier()
```

```
In [31]: pred=dtc.predict(x_test)
```

```
In [32]: pred
```

```
Out[32]: array([10, 25, 18, 20, 24,  3, 24,  2,  1, 14, 20, 23,  2,  1, 20,  2,  3,
                3, 10, 15,  9,  6,  3, 15, 23,  8, 25,  2,  5,  2,  2,  2, 18,  7,
               10,  3, 24, 12,  9, 18, 29, 13, 10,  3,  1,  5,  2,  2, 10,  1, 15,
               20,  9,  1,  2,  3,  4,  8,  8,  1,  2, 18,  2,  1, 26,  3,  2, 29,
               9,  9, 25, 23, 20, 19,  4, 19,  9,  2,  2,  2,  4,  8,  9,  1, 20,
               5, 14,  1,  3, 29, 14, 27, 16,  7, 10, 16, 25,  2, 10,  2, 13, 29,
               22, 14, 10,  2,  6, 10, 29,  1, 16,  6,  1, 10,  9, 10,  1,  6,  4,
               3, 29,  2,  5, 24,  1,  2,  2,  8,  7,  5, 12,  9,  2,  7,  7,  6,
               5,  6,  7,  2,  6, 10,  8,  1,  6,  1,  2,  4,  3,  3,  1, 17,  1,
               2, 18,  8, 29,  2, 21,  1, 12,  1,  9, 25,  1, 17,  1, 18, 20,  3,
               26,  2,  8, 11, 12,  1, 10,  1,  7,  9, 20,  9,  2, 24, 11,  2, 16,
               1, 23,  2, 28,  8, 15,  1,  7,  2,  2, 17, 12,  2,  7,  8,  9,  5,
               27, 20,  1,  1,  7,  3, 10, 28,  8,  1,  8, 10,  5, 10,  9,  1, 12,
               2, 26,  8,  3, 13,  8,  2,  9,  6,  7,  2,  1, 10,  6,  8,  4, 18,
               19, 23,  1,  4,  2,  9,  1, 10, 11,  2,  7,  7, 19,  2, 29, 14,  4,
               7, 18, 10, 16,  9, 10,  1, 22,  7,  2,  9,  8,  7, 13,  6,  1,  7,
               26, 16,  7, 11, 18,  2, 19,  6,  4,  1,  4,  2,  1,  6,  2, 22,  3,
               3,  1,  6,  2,  8], dtype=int64)
```

```
In [33]: y_test
```


```
Out[33]: 442      10
          1091    25
          981    18
          785    20
          1332   24
          ..
          1439    3
          481     1
          124     6
          198     2
          1229    8
          Name: DistanceFromHome, Length: 294, dtype: int64
```

```
In [34]: df
```

Out[34]:

	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Ed
0	41	Yes	Travel_Rarely	1102	Sales		1
1	49	No	Travel_Frequently	279	Research & Development		8
2	37	Yes	Travel_Rarely	1373	Research & Development		2
3	33	No	Travel_Frequently	1392	Research & Development		3
4	27	No	Travel_Rarely	591	Research & Development		2
...
1465	36	No	Travel_Frequently	884	Research & Development		23
1466	39	No	Travel_Rarely	613	Research & Development		6
1467	27	No	Travel_Rarely	155	Research & Development		4
1468	49	No	Travel_Frequently	1023	Sales		2
1469	34	No	Travel_Rarely	628	Research & Development		8

1470 rows × 35 columns

In [35]: `dtc.predict(ms.transform([[1,20,18000]]))`

C:\Users\chatu\anaconda3\Lib\site-packages\sklearn\base.py:464: UserWarning: X does not have valid feature names, but MinMaxScaler was fitted with feature names
 warnings.warn(
 C:\Users\chatu\anaconda3\Lib\site-packages\sklearn\base.py:464: UserWarning: X does not have valid feature names, but DecisionTreeClassifier was fitted with feature names
 warnings.warn(

Out[35]: `array([20], dtype=int64)`

Evaluation of classification model

In [36]: `#Accuracy score
 from sklearn.metrics import accuracy_score, confusion_matrix, classification_report`

In [37]: `accuracy_score(y_test, pred)`Out[37]: `1.0`In [38]: `confusion_matrix(y_test, pred)`

[illegible]


```
In [39]: pd.crosstab(y_test, pred)
```

Out[39]:

col_012345678910...20212223242

DistanceFromHome																												
1	37	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0											
2	0	45	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0											
3	0	0	17	0	0	0	0	0	0	0	0	...	0	0	0	0	0											
4	0	0	0	10	0	0	0	0	0	0	0	...	0	0	0	0	0											
5	0	0	0	0	8	0	0	0	0	0	0	...	0	0	0	0	0											
6	0	0	0	0	0	14	0	0	0	0	0	...	0	0	0	0	0											
7	0	0	0	0	0	0	18	0	0	0	0	...	0	0	0	0	0											
8	0	0	0	0	0	0	0	17	0	0	0	...	0	0	0	0	0											
9	0	0	0	0	0	0	0	0	18	0	0	...	0	0	0	0	0											
10	0	0	0	0	0	0	0	0	0	0	20	...	0	0	0	0	0											
11	0	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0											
12	0	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0											
13	0	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0											
14	0	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0											
15	0	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0											
16	0	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0											
17	0	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0											
18	0	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0											
19	0	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0											
20	0	0	0	0	0	0	0	0	0	0	0	...	9	0	0	0	0											
21	0	0	0	0	0	0	0	0	0	0	0	...	0	1	0	0	0											
22	0	0	0	0	0	0	0	0	0	0	0	...	0	0	3	0	0											
23	0	0	0	0	0	0	0	0	0	0	0	...	0	0	0	5	0											
24	0	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	5											
25	0	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0											
26	0	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0											
27	0	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0											
28	0	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0											
29	0	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0											

29 rows × 29 columns



predicted no predicted yes

Actual No 58=TN 0=FP Actual yes 6=FN 16=TP

In [41]: `(58+16)/80 #accuracy`

Out[41]: 0.925

In [42]: `print(classification_report(y_test,pred))`

	precision	recall	f1-score	support
1	1.00	1.00	1.00	37
2	1.00	1.00	1.00	45
3	1.00	1.00	1.00	17
4	1.00	1.00	1.00	10
5	1.00	1.00	1.00	8
6	1.00	1.00	1.00	14
7	1.00	1.00	1.00	18
8	1.00	1.00	1.00	17
9	1.00	1.00	1.00	18
10	1.00	1.00	1.00	20
11	1.00	1.00	1.00	4
12	1.00	1.00	1.00	6
13	1.00	1.00	1.00	4
14	1.00	1.00	1.00	5
15	1.00	1.00	1.00	4
16	1.00	1.00	1.00	6
17	1.00	1.00	1.00	3
18	1.00	1.00	1.00	9
19	1.00	1.00	1.00	5
20	1.00	1.00	1.00	9
21	1.00	1.00	1.00	1
22	1.00	1.00	1.00	3
23	1.00	1.00	1.00	5
24	1.00	1.00	1.00	5
25	1.00	1.00	1.00	5
26	1.00	1.00	1.00	4
27	1.00	1.00	1.00	2
28	1.00	1.00	1.00	2
29	1.00	1.00	1.00	8
accuracy				1.00 294
macro avg				1.00 294
weighted avg				1.00 294

Roc-AUC curve

In [43]: `probability=dtc.predict_proba(x_test)[:,-1]`

In [44]: `probability`

```
Out[44]: array([0., 0., 0., 0., 0., 0., 0., 1., 0., 0., 0., 0., 1., 0., 0., 1., 0.,
                0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 1., 0., 1., 1., 1., 0., 0.,
                0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 1., 1., 0., 0., 0.,
                0., 0., 0., 1., 0., 0., 0., 0., 0., 1., 0., 1., 0., 0., 0., 1., 0.,
                0., 0., 0., 0., 0., 0., 0., 0., 0., 1., 1., 1., 0., 0., 0., 0., 0.,
                0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 1., 0., 1., 0., 0.,
                0., 0., 0., 1., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0.,
                0., 0., 1., 0., 0., 0., 1., 1., 0., 0., 0., 0., 0., 1., 0., 0., 0.,
                0., 0., 0., 1., 0., 0., 0., 0., 0., 0., 1., 0., 0., 0., 0., 0.,
                1., 0., 0., 0., 1., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0.,
                0., 1., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 1., 0., 0., 1., 0.,
                0., 0., 1., 0., 0., 0., 0., 0., 1., 1., 0., 0., 1., 0., 0., 0.,
                0., 0., 0., 0., 0., 0., 1., 0., 0., 0., 1., 0., 0., 0., 0., 0.,
                0., 0., 0., 0., 1., 0., 0., 0., 0., 1., 0., 0., 0., 1., 0., 0.,
                0., 0., 0., 0., 0., 0., 0., 0., 0., 1., 0., 0., 0., 0., 0., 0.,
                0., 0., 0., 0., 0., 1., 0., 0., 0., 0., 0., 1., 0., 0., 1., 0., 0.,
                0., 0., 0., 1., 0.]
```

```
In [45]: from sklearn import tree
plt.figure(figsize=(25,15))
tree.plot_tree(dtc,filled=True)
```

```

Out[45]: [Text(0.1095890410958904, 0.9722222222222222, 'x[1] <= 0.018\ngini = 0.932\nsam
ples = 1176\nvalue = [171, 166, 67, 54, 57, 45, 66, 63, 67, 66, 25, 14\n15, 16,
22, 26, 17, 17, 17, 16, 17, 16, 22, 23\n20, 21, 10, 21, 19]'),
Text(0.0547945205479452, 0.9166666666666666, 'gini = 0.0\nsamples = 171\nvalue
= [171, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0\n0]'),
Text(0.1643835616438356, 0.9166666666666666, 'x[1] <= 0.054\ngini = 0.936\nsam
ples = 1005\nvalue = [0, 166, 67, 54, 57, 45, 66, 63, 67, 66, 25, 14\n15, 16, 2
2, 26, 17, 17, 17, 16, 17, 16, 22, 23\n20, 21, 10, 21, 19]'),
Text(0.1095890410958904, 0.8611111111111112, 'gini = 0.0\nsamples = 166\nvalue
= [0, 166, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0\n0]'),
Text(0.2191780821917808, 0.8611111111111112, 'x[1] <= 0.089\ngini = 0.948\nsam
ples = 839\nvalue = [0, 0, 67, 54, 57, 45, 66, 63, 67, 66, 25, 14\n15, 16, 22,
26, 17, 17, 17, 16, 17, 16, 22, 23\n20, 21, 10, 21, 19]'),
Text(0.1643835616438356, 0.8055555555555556, 'gini = 0.0\nsamples = 67\nvalue
= [0, 0, 67, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0\n0]'),
Text(0.273972602739726, 0.8055555555555556, 'x[1] <= 0.161\ngini = 0.946\nsamp
les = 772\nvalue = [0, 0, 0, 54, 57, 45, 66, 63, 67, 66, 25, 14\n15, 16, 22, 2
6, 17, 17, 17, 16, 17, 16, 22, 23\n20, 21, 10, 21, 19]'),
Text(0.136986301369863, 0.75, 'x[1] <= 0.125\ngini = 0.5\nsamples = 111\nvalue
= [0, 0, 0, 54, 57, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0\n0]'),
Text(0.0821917808219178, 0.6944444444444444, 'gini = 0.0\nsamples = 54\nvalue
= [0, 0, 0, 54, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0\n0]'),
Text(0.1917808219178082, 0.6944444444444444, 'gini = 0.0\nsamples = 57\nvalue
= [0, 0, 0, 0, 57, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0\n0]'),
Text(0.410958904109589, 0.75, 'x[1] <= 0.232\ngini = 0.94\nsamples = 661\nvalu
e = [0, 0, 0, 0, 0, 45, 66, 63, 67, 66, 25, 14, 15\n16, 22, 26, 17, 17, 17, 16,
17, 16, 22, 23, 20\n21, 10, 21, 19]'),
Text(0.3013698630136986, 0.6944444444444444, 'x[1] <= 0.196\ngini = 0.482\nsam
ples = 111\nvalue = [0, 0, 0, 0, 0, 45, 66, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0\n0]'),
Text(0.2465753424657534, 0.6388888888888888, 'gini = 0.0\nsamples = 45\nvalue
= [0, 0, 0, 0, 0, 45, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0\n0]'),
Text(0.3561643835616438, 0.6388888888888888, 'gini = 0.0\nsamples = 66\nvalue
= [0, 0, 0, 0, 0, 0, 66, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0\n0]'),
Text(0.5205479452054794, 0.6944444444444444, 'x[1] <= 0.268\ngini = 0.935\nsam
ples = 550\nvalue = [0, 0, 0, 0, 0, 0, 0, 63, 67, 66, 25, 14, 15\n16, 22, 26, 1
7, 17, 17, 16, 17, 16, 22, 23, 20\n21, 10, 21, 19]'),
Text(0.4657534246575342, 0.6388888888888888, 'gini = 0.0\nsamples = 63\nvalue
= [0, 0, 0, 0, 0, 0, 0, 63, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0\n0]'),
Text(0.5753424657534246, 0.6388888888888888, 'x[1] <= 0.304\ngini = 0.934\nsam
ples = 487\nvalue = [0, 0, 0, 0, 0, 0, 0, 0, 67, 66, 25, 14, 15\n16, 22, 26, 1
7, 17, 17, 16, 17, 16, 22, 23, 20\n21, 10, 21, 19]'),
Text(0.5205479452054794, 0.5833333333333334, 'gini = 0.0\nsamples = 67\nvalue
= [0, 0, 0, 0, 0, 0, 0, 67, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0\n0]'),
Text(0.6301369863013698, 0.5833333333333334, 'x[1] <= 0.339\ngini = 0.936\nsam
ples = 420\nvalue = [0, 0, 0, 0, 0, 0, 0, 0, 66, 25, 14, 15\n16, 22, 26, 17,
17, 17, 16, 17, 16, 22, 23, 20\n21, 10, 21, 19]'),
Text(0.5753424657534246, 0.5277777777777778, 'gini = 0.0\nsamples = 66\nvalue
= [0, 0, 0, 0, 0, 0, 0, 0, 66, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0\n0]'),

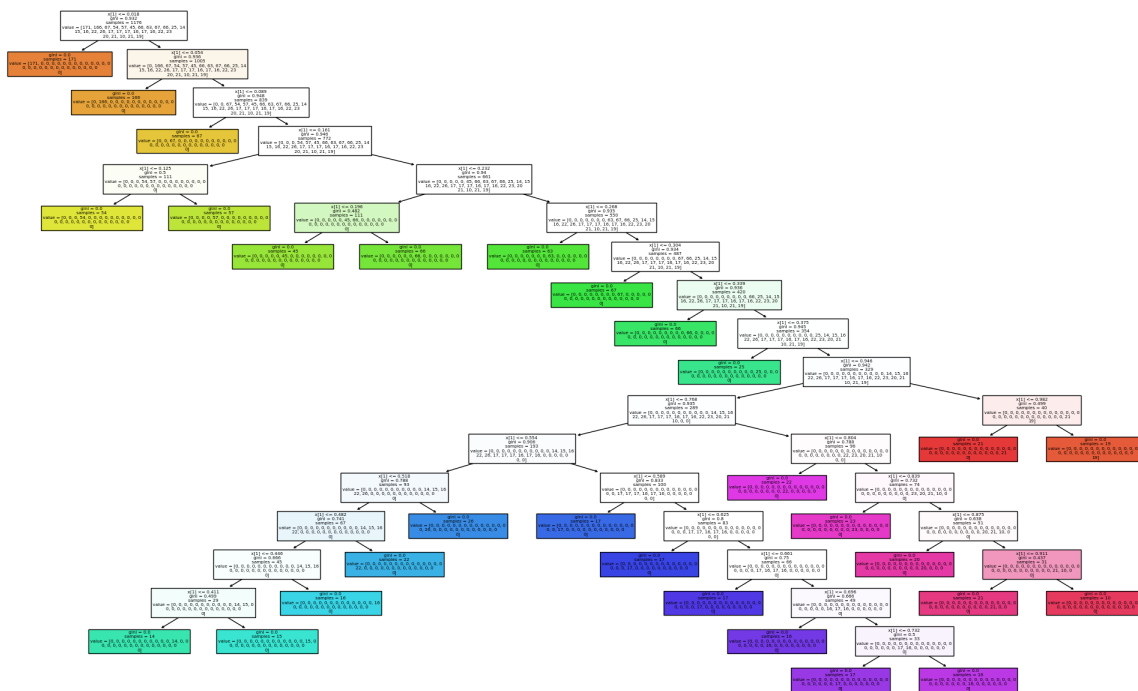
```

22/35

```

17, 16, 0, 0, 0, 0, 0, 0, 0\n0']'),
Text(0.6712328767123288, 0.08333333333333333, 'gini = 0.0\nsamples = 16\nvalue
= [0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 16, 0, 0, 0, 0, 0,
0, 0, 0\n0']'),
Text(0.7808219178082192, 0.08333333333333333, 'x[1] <= 0.732\ngini = 0.5\nsampl
es = 33\nvalue = [0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
17, 16, 0, 0, 0, 0, 0, 0, 0, 0\n0']'),
Text(0.726027397260274, 0.027777777777777776, 'gini = 0.0\nsamples = 17\nvalue
= [0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 17, 0, 0, 0, 0,
0, 0, 0\n0']'),
Text(0.8356164383561644, 0.027777777777777776, 'gini = 0.0\nsamples = 16\nvalu
e = [0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 16, 0, 0,
0, 0, 0, 0\n0']'),
Text(0.726027397260274, 0.3611111111111111, 'x[1] <= 0.804\ngini = 0.788\nsamp
les = 96\nvalue = [0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 22, 23, 20, 21, 10\n0, 0]'),
Text(0.6712328767123288, 0.30555555555555556, 'gini = 0.0\nsamples = 22\nvalue
= [0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 22, 0, 0,
0, 0, 0\n0']'),
Text(0.7808219178082192, 0.30555555555555556, 'x[1] <= 0.839\ngini = 0.732\nsam
ples = 74\nvalue = [0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 23, 20, 21, 10, 0\n0]'),
Text(0.726027397260274, 0.25, 'gini = 0.0\nsamples = 23\nvalue = [0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 23, 0, 0, 0, 0, 0\n0]'),
Text(0.8356164383561644, 0.25, 'x[1] <= 0.875\ngini = 0.638\nsamples = 51\nval
ue = [0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 21, 10, 0\n0]'),
Text(0.7808219178082192, 0.19444444444444445, 'gini = 0.0\nsamples = 20\nvalue
= [0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 20,
0, 0, 0\n0]'),
Text(0.8904109589041096, 0.19444444444444445, 'x[1] <= 0.911\ngini = 0.437\nsa
mples = 31\nvalue = [0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 21, 10, 0\n0]'),
Text(0.8356164383561644, 0.13888888888888889, 'gini = 0.0\nsamples = 21\nvalue
= [0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 2
1, 0, 0\n0]'),
Text(0.9452054794520548, 0.13888888888888889, 'gini = 0.0\nsamples = 10\nvalue
= [0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 10, 0\n0]'),
Text(0.8904109589041096, 0.41666666666666667, 'x[1] <= 0.982\ngini = 0.499\nsam
ples = 40\nvalue = [0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 21\n19]'),
Text(0.8356164383561644, 0.3611111111111111, 'gini = 0.0\nsamples = 21\nvalue
= [0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 21\n0]'),
Text(0.9452054794520548, 0.3611111111111111, 'gini = 0.0\nsamples = 19\nvalue
= [0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0\n19]')]]

```



```
In [46]: from sklearn.model_selection import GridSearchCV
parameter={
    'criterion':['gini','entropy'],
    'splitter':['best','random'],
    'max_depth':[1,2,3,4,5],
    'max_features':['auto', 'sqrt', 'log2']
}
```

```
In [47]: grid_search=GridSearchCV(estimator=dtc,param_grid=parameter,cv=5,scoring="accuracy")
```

```
In [48]: grid_search.fit(x_train,y_train)
```



```
C:\Users\chatu\anaconda3\Lib\site-packages\sklearn\model_selection\_validation.py:425: FitFailedWarning:
100 fits failed out of a total of 300.
The score on these train-test partitions for these parameters will be set to nan.
If these failures are not expected, you can try to debug them by setting error_score='raise'.
```

Below are more details about the failures:

```
-----
100 fits failed with the following error:
Traceback (most recent call last):
  File "C:\Users\chatu\anaconda3\Lib\site-packages\sklearn\model_selection\_validation.py", line 732, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
  File "C:\Users\chatu\anaconda3\Lib\site-packages\sklearn\base.py", line 1144, in wrapper
    estimator._validate_params()
  File "C:\Users\chatu\anaconda3\Lib\site-packages\sklearn\base.py", line 637, in _validate_params
    validate_parameter_constraints(
  File "C:\Users\chatu\anaconda3\Lib\site-packages\sklearn\utils\_param_validation.py", line 95, in validate_parameter_constraints
    raise InvalidParameterError(
sklearn.utils._param_validation.InvalidParameterError: The 'max_features' parameter of DecisionTreeClassifier must be an int in the range [1, inf), a float in the range (0.0, 1.0], a str among {'sqrt', 'log2'} or None. Got 'auto' instead.
```

```
warnings.warn(some_fits_failed_message, FitFailedWarning)
C:\Users\chatu\anaconda3\Lib\site-packages\sklearn\model_selection\_search.py:976: UserWarning: One or more of the test scores are non-finite: [      nan
nan 0.23390912 0.1641291 0.26443924 0.16072845
      nan      nan 0.19476019 0.16327804 0.22870537 0.17773891
      nan      nan 0.2041291 0.1981392 0.29338262 0.19558601
      nan      nan 0.30954922 0.16496574 0.26795889 0.26097367
      nan      nan 0.33679409 0.29338983 0.30106022 0.25424089
      nan      nan 0.17860079 0.16156149 0.17860079 0.16326722
      nan      nan 0.15561125 0.15393076 0.17179228 0.20485034
      nan      nan 0.20491165 0.16751893 0.25003967 0.17263613
      nan      nan 0.26614497 0.24069239 0.24834115 0.24322395
      nan      nan 0.31804904 0.23473134 0.29600793 0.22699243]
warnings.warn(
```

```
Out[48]: ▶ GridSearchCV
          ▶ estimator: DecisionTreeClassifier
            ▶ DecisionTreeClassifier
```

```
In [49]: dtc_cv=DecisionTreeClassifier(criterion='entropy',
    max_depth=3,
    max_features='sqrt',
    splitter='best')
    dtc_cv.fit(x_train,y_train)
```

Out[49]:

DecisionTreeClassifier

```
DecisionTreeClassifier(criterion='entropy', max_depth=3, max_features='sqrt')
```

In [50]: `pred=dtc_cv.predict(x_test)`In [51]: `print(classification_report(y_test,pred))`

	precision	recall	f1-score	support
1	0.30	0.86	0.45	37
2	0.40	0.22	0.29	45
3	0.00	0.00	0.00	17
4	0.00	0.00	0.00	10
5	0.00	0.00	0.00	8
6	0.00	0.00	0.00	14
7	0.21	0.94	0.35	18
8	0.00	0.00	0.00	17
9	0.25	0.06	0.09	18
10	0.00	0.00	0.00	20
11	0.00	0.00	0.00	4
12	0.00	0.00	0.00	6
13	0.00	0.00	0.00	4
14	0.00	0.00	0.00	5
15	0.00	0.00	0.00	4
16	0.08	1.00	0.15	6
17	0.00	0.00	0.00	3
18	0.00	0.00	0.00	9
19	0.00	0.00	0.00	5
20	0.00	0.00	0.00	9
21	0.00	0.00	0.00	1
22	0.00	0.00	0.00	3
23	0.00	0.00	0.00	5
24	0.00	0.00	0.00	5
25	0.00	0.00	0.00	5
26	0.00	0.00	0.00	4
27	0.00	0.00	0.00	2
28	0.00	0.00	0.00	2
29	0.00	0.00	0.00	8
accuracy			0.22	294
macro avg	0.04	0.11	0.05	294
weighted avg	0.13	0.22	0.13	294

```
C:\Users\chatu\anaconda3\Lib\site-packages\sklearn\metrics\_classification.py:146
9: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to
0.0 in labels with no predicted samples. Use `zero_division` parameter to control
this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
C:\Users\chatu\anaconda3\Lib\site-packages\sklearn\metrics\_classification.py:146
9: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to
0.0 in labels with no predicted samples. Use `zero_division` parameter to control
this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
C:\Users\chatu\anaconda3\Lib\site-packages\sklearn\metrics\_classification.py:146
9: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to
0.0 in labels with no predicted samples. Use `zero_division` parameter to control
this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
```

Random Forest

```
In [52]: from sklearn.ensemble import RandomForestClassifier
rfc=RandomForestClassifier()

In [53]: forest_params = [{'max_depth': list(range(10, 15)), 'max_features': list(range(0.5, 1.0))}]

In [54]: rfc_cv= GridSearchCV(rfc,param_grid=forest_params,cv=10,scoring="accuracy")

In [55]: # Fit the GridSearchCV to your training data
rfc_cv.fit(x_train, y_train)
```

```

-----
KeyboardInterrupt                                Traceback (most recent call last)
Cell In[55], line 2
      1 # Fit the GridSearchCV to your training data
----> 2 rfc_cv.fit(x_train, y_train)

File ~\anaconda3\Lib\site-packages\sklearn\base.py:1151, in _fit_context.<locals>
>.decorator.<locals>.wrapper(estimator, *args, **kwargs)
    1144     estimator._validate_params()
    1146     with config_context(
    1147         skip_parameter_validation=(
    1148             prefer_skip_nested_validation or global_skip_validation
    1149         )
    1150     ):
-> 1151         return fit_method(estimator, *args, **kwargs)

File ~\anaconda3\Lib\site-packages\sklearn\model_selection\_search.py:898, in Base
eSearchCV.fit(self, X, y, groups, **fit_params)
    892     results = self._format_results(
    893         all_candidate_params, n_splits, all_out, all_more_results
    894     )
    896     return results
--> 898 self._run_search(evaluate_candidates)
    900 # multimetric is determined here because in the case of a callable
    901 # self.scoring the return type is only known after calling
    902 first_test_score = all_out[0]["test_scores"]

File ~\anaconda3\Lib\site-packages\sklearn\model_selection\_search.py:1419, in Gr
idSearchCV._run_search(self, evaluate_candidates)
    1417 def _run_search(self, evaluate_candidates):
    1418     """Search all candidates in param_grid"""
-> 1419     evaluate_candidates(ParameterGrid(self.param_grid))

File ~\anaconda3\Lib\site-packages\sklearn\model_selection\_search.py:845, in Bas
eSearchCV.fit.<locals>.evaluate_candidates(candidate_params, cv, more_results)
    837 if self.verbose > 0:
    838     print(
    839         "Fitting {0} folds for each of {1} candidates,"
    840         " totalling {2} fits".format(
    841             n_splits, n_candidates, n_candidates * n_splits
    842         )
    843     )
--> 845 out = parallel(
    846     delayed(_fit_and_score)(
    847         clone(base_estimator),
    848         X,
    849         y,
    850         train=train,
    851         test=test,
    852         parameters=parameters,
    853         split_progress=(split_idx, n_splits),
    854         candidate_progress=(cand_idx, n_candidates),
    855         **fit_and_score_kwargs,
    856     )
    857     for (cand_idx, parameters), (split_idx, (train, test)) in product(
    858         enumerate(candidate_params), enumerate(cv.split(X, y, groups))
    859     )
    860 )
    862 if len(out) < 1:
    863     raise ValueError(

```

```

864         "No fits were performed. "
865         "Was the CV iterator empty? "
866         "Were there no candidates?"
867     )

```

File ~\anaconda3\Lib\site-packages\sklearn\utils\parallel.py:65, in Parallel.__call__(self, iterable)

```

60 config = get_config()
61 iterable_with_config = (
62     (_with_config(delayed_func, config), args, kwargs)
63     for delayed_func, args, kwargs in iterable
64 )
--> 65 return super().__call__(iterable_with_config)

```

File ~\anaconda3\Lib\site-packages\joblib\parallel.py:1088, in Parallel.__call__(self, iterable)

```

1085 if self.dispatch_one_batch(iterator):
1086     self._iterating = self._original_iterator is not None
-> 1088 while self.dispatch_one_batch(iterator):
1089     pass
1091 if pre_dispatch == "all" or n_jobs == 1:
1092     # The iterable was consumed all at once by the above for loop.
1093     # No need to wait for async callbacks to trigger to
1094     # consumption.

```

File ~\anaconda3\Lib\site-packages\joblib\parallel.py:901, in Parallel.dispatch_one_batch(self, iterator)

```

899     return False
900 else:
--> 901     self._dispatch(tasks)
902     return True

```

File ~\anaconda3\Lib\site-packages\joblib\parallel.py:819, in Parallel._dispatch(self, batch)

```

817 with self._lock:
818     job_idx = len(self._jobs)
--> 819     job = self._backend.apply_async(batch, callback=cb)
820     # A job can complete so quickly than its callback is
821     # called before we get here, causing self._jobs to
822     # grow. To ensure correct results ordering, .insert is
823     # used (rather than .append) in the following line
824     self._jobs.insert(job_idx, job)

```

File ~\anaconda3\Lib\site-packages\joblib_parallel_backends.py:208, in SequentialBackend.apply_async(self, func, callback)

```

206 def apply_async(self, func, callback=None):
207     """Schedule a func to be run"""
--> 208     result = ImmediateResult(func)
209     if callback:
210         callback(result)

```

File ~\anaconda3\Lib\site-packages\joblib_parallel_backends.py:597, in ImmediateResult.__init__(self, batch)

```

594 def __init__(self, batch):
595     # Don't delay the application, to avoid keeping the input
596     # arguments in memory
--> 597     self.results = batch()

```

File ~\anaconda3\Lib\site-packages\joblib\parallel.py:288, in BatchedCalls.__call__(self)

```

284 def __call__(self):
285     # Set the default nested backend to self._backend but do not set the
286     # change the default number of processes to -1
287     with parallel_backend(self._backend, n_jobs=self._n_jobs):
--> 288         return [func(*args, **kwargs)
289                 for func, args, kwargs in self.items]

File ~\anaconda3\Lib\site-packages\joblib\parallel.py:288, in <listcomp>(.0)
284 def __call__(self):
285     # Set the default nested backend to self._backend but do not set the
286     # change the default number of processes to -1
287     with parallel_backend(self._backend, n_jobs=self._n_jobs):
--> 288         return [func(*args, **kwargs)
289                 for func, args, kwargs in self.items]

File ~\anaconda3\Lib\site-packages\sklearn\utils\parallel.py:127, in _FuncWrapper.__call__(self, *args, **kwargs)
125     config = {}
126 with config_context(**config):
--> 127     return self.function(*args, **kwargs)

File ~\anaconda3\Lib\site-packages\sklearn\model_selection\_validation.py:732, in _fit_and_score(estimator, X, y, scorer, train, test, verbose, parameters, fit_params, return_train_score, return_parameters, return_n_test_samples, return_times, return_estimator, split_progress, candidate_progress, error_score)
730     estimator.fit(X_train, **fit_params)
731 else:
--> 732     estimator.fit(X_train, y_train, **fit_params)
734 except Exception:
735     # Note fit time as time until error
736     fit_time = time.time() - start_time

File ~\anaconda3\Lib\site-packages\sklearn\base.py:1151, in _fit_context.<locals>.<decorator.<locals>.wrapper(estimator, *args, **kwargs)
1144     estimator._validate_params()
1146 with config_context(
1147     skip_parameter_validation=(
1148         prefer_skip_nested_validation or global_skip_validation
1149     )
1150 ):
-> 1151     return fit_method(estimator, *args, **kwargs)

File ~\anaconda3\Lib\site-packages\sklearn\ensemble\_forest.py:456, in BaseForest.fit(self, X, y, sample_weight)
445 trees = [
446     self._make_estimator(append=False, random_state=random_state)
447     for i in range(n_more_estimators)
448 ]
450 # Parallel loop: we prefer the threading backend as the Cython code
451 # for fitting the trees is internally releasing the Python GIL
452 # making threading more efficient than multiprocessing in
453 # that case. However, for joblib 0.12+ we respect any
454 # parallel_backend contexts set at a higher level,
455 # since correctness does not rely on using threads.
--> 456 trees = Parallel(
457     n_jobs=self.n_jobs,
458     verbose=self.verbose,
459     prefer="threads",
460 )(
461     delayed(_parallel_build_trees)(

```

```

462         t,
463         self.bootstrap,
464         X,
465         y,
466         sample_weight,
467         i,
468         len(trees),
469         verbose=self.verbose,
470         class_weight=self.class_weight,
471         n_samples_bootstrap=n_samples_bootstrap,
472     )
473     for i, t in enumerate(trees)
474 )
475 # Collect newly grown trees
476 self.estimateds_.extend(trees)

```

File ~\anaconda3\Lib\site-packages\sklearn\utils\parallel.py:65, in Parallel.__call__(self, iterable)

```

60 config = get_config()
61 iterable_with_config = (
62     (_with_config(delayed_func, config), args, kwargs)
63     for delayed_func, args, kwargs in iterable
64 )
--> 65 return super().__call__(iterable_with_config)

```

File ~\anaconda3\Lib\site-packages\joblib\parallel.py:1088, in Parallel.__call__(self, iterable)

```

1085 if self.dispatch_one_batch(iterator):
1086     self._iterating = self._original_iterator is not None
-> 1088 while self.dispatch_one_batch(iterator):
1089     pass
1091 if pre_dispatch == "all" or n_jobs == 1:
1092     # The iterable was consumed all at once by the above for loop.
1093     # No need to wait for async callbacks to trigger to
1094     # consumption.

```

File ~\anaconda3\Lib\site-packages\joblib\parallel.py:901, in Parallel.dispatch_one_batch(self, iterator)

```

899     return False
900 else:
--> 901     self._dispatch(tasks)
902     return True

```

File ~\anaconda3\Lib\site-packages\joblib\parallel.py:819, in Parallel._dispatch(self, batch)

```

817 with self._lock:
818     job_idx = len(self._jobs)
--> 819     job = self._backend.apply_async(batch, callback=cb)
820     # A job can complete so quickly than its callback is
821     # called before we get here, causing self._jobs to
822     # grow. To ensure correct results ordering, .insert is
823     # used (rather than .append) in the following line
824     self._jobs.insert(job_idx, job)

```

File ~\anaconda3\Lib\site-packages\joblib_parallel_backends.py:208, in SequentialBackend.apply_async(self, func, callback)

```

206 def apply_async(self, func, callback=None):
207     """Schedule a func to be run"""
--> 208     result = ImmediateResult(func)
209     if callback:

```

```

210         callback(result)

File ~\anaconda3\Lib\site-packages\joblib\parallel_backends.py:597, in ImmediateResult.__init__(self, batch)
    594 def __init__(self, batch):
    595     # Don't delay the application, to avoid keeping the input
    596     # arguments in memory
--> 597     self.results = batch()

File ~\anaconda3\Lib\site-packages\joblib\parallel.py:288, in BatchedCalls.__call__(self)
    284 def __call__(self):
    285     # Set the default nested backend to self._backend but do not set the
    286     # change the default number of processes to -1
    287     with parallel_backend(self._backend, n_jobs=self._n_jobs):
--> 288         return [func(*args, **kwargs)
    289                 for func, args, kwargs in self.items]

File ~\anaconda3\Lib\site-packages\joblib\parallel.py:288, in <listcomp>(.0)
    284 def __call__(self):
    285     # Set the default nested backend to self._backend but do not set the
    286     # change the default number of processes to -1
    287     with parallel_backend(self._backend, n_jobs=self._n_jobs):
--> 288         return [func(*args, **kwargs)
    289                 for func, args, kwargs in self.items]

File ~\anaconda3\Lib\site-packages\sklearn\utils\parallel.py:127, in _FuncWrapper.__call__(self, *args, **kwargs)
    125     config = {}
    126     with config_context(**config):
--> 127         return self.function(*args, **kwargs)

File ~\anaconda3\Lib\site-packages\sklearn\ensemble\_forest.py:188, in _parallel_build_trees(tree, bootstrap, X, y, sample_weight, tree_idx, n_trees, verbose, class_weight, n_samples_bootstrap)
    185     elif class_weight == "balanced_subsample":
    186         curr_sample_weight *= compute_sample_weight("balanced", y, indices=indices)
--> 188     tree.fit(X, y, sample_weight=curr_sample_weight, check_input=False)
    189 else:
    190     tree.fit(X, y, sample_weight=sample_weight, check_input=False)

File ~\anaconda3\Lib\site-packages\sklearn\base.py:1151, in _fit_context.<locals>.<decorator.<locals>.wrapper(estimator, *args, **kwargs)
    1144     estimator._validate_params()
    1146     with config_context(
    1147         skip_parameter_validation=(
    1148             prefer_skip_nested_validation or global_skip_validation
    1149         )
    1150     ):
-> 1151         return fit_method(estimator, *args, **kwargs)

File ~\anaconda3\Lib\site-packages\sklearn\tree\_classes.py:959, in DecisionTreeClassifier.fit(self, X, y, sample_weight, check_input)
    928 @_fit_context(prefer_skip_nested_validation=True)
    929 def fit(self, X, y, sample_weight=None, check_input=True):
    930     """Build a decision tree classifier from the training set (X, y).
    931
    932     Parameters
    933     (...)

```



```

956         Fitted estimator.
957         """
--> 959     super()._fit(
960         X,
961         y,
962         sample_weight=sample_weight,
963         check_input=check_input,
964     )
965     return self

File ~\anaconda3\Lib\site-packages\sklearn\tree\_classes.py:443, in BaseDecisionTree._fit(self, X, y, sample_weight, check_input, missing_values_in_feature_mask)
432 else:
433     builder = BestFirstTreeBuilder(
434         splitter,
435         min_samples_split,
436         (...)
437         self.min_impurity_decrease,
438     )
--> 443 builder.build(self.tree_, X, y, sample_weight, missing_values_in_feature_mask)
444 if self.n_outputs_ == 1 and is_classifier(self):
445     self.n_classes_ = self.n_classes_[0]

KeyboardInterrupt:

```

```
In [ ]: pred = rfc_cv.predict(x_test)
```

```
In [56]: print(classification_report(y_test, pred))
```

	precision	recall	f1-score	support
1	0.30	0.86	0.45	37
2	0.40	0.22	0.29	45
3	0.00	0.00	0.00	17
4	0.00	0.00	0.00	10
5	0.00	0.00	0.00	8
6	0.00	0.00	0.00	14
7	0.21	0.94	0.35	18
8	0.00	0.00	0.00	17
9	0.25	0.06	0.09	18
10	0.00	0.00	0.00	20
11	0.00	0.00	0.00	4
12	0.00	0.00	0.00	6
13	0.00	0.00	0.00	4
14	0.00	0.00	0.00	5
15	0.00	0.00	0.00	4
16	0.08	1.00	0.15	6
17	0.00	0.00	0.00	3
18	0.00	0.00	0.00	9
19	0.00	0.00	0.00	5
20	0.00	0.00	0.00	9
21	0.00	0.00	0.00	1
22	0.00	0.00	0.00	3
23	0.00	0.00	0.00	5
24	0.00	0.00	0.00	5
25	0.00	0.00	0.00	5
26	0.00	0.00	0.00	4
27	0.00	0.00	0.00	2
28	0.00	0.00	0.00	2
29	0.00	0.00	0.00	8
accuracy				0.22
macro avg				0.04
weighted avg				0.13

C:\Users\chatu\anaconda3\Lib\site-packages\sklearn\metrics_classification.py:146
 9: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to
 0.0 in labels with no predicted samples. Use `zero_division` parameter to control
 this behavior.

_warn_prf(average, modifier, msg_start, len(result))

C:\Users\chatu\anaconda3\Lib\site-packages\sklearn\metrics_classification.py:146
 9: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to
 0.0 in labels with no predicted samples. Use `zero_division` parameter to control
 this behavior.

_warn_prf(average, modifier, msg_start, len(result))

C:\Users\chatu\anaconda3\Lib\site-packages\sklearn\metrics_classification.py:146
 9: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to
 0.0 in labels with no predicted samples. Use `zero_division` parameter to control
 this behavior.

_warn_prf(average, modifier, msg_start, len(result))

```
In [58]: from sklearn.ensemble import RandomForestClassifier
rf_model = RandomForestClassifier()
rf_model.fit(x_train, y_train)
```

Out[58]: ▾ RandomForestClassifier
RandomForestClassifier()